#klimatapotrebuje

Klima Kompas Platform

Komunálne Voľby 2026

Initiative Document — Version 5.0 (LOCKED)

Version 5.0 | April 2026 | Methodology locked — implementation ready

Confidential — Internal Use Only

1. PROBLEM & CONTEXT

Slovak voters in komunálne voľby 2026 have no structured, neutral source for evaluating how candidates for regional and city leadership approach climate and environmental policy. No tool exists that takes publicly available evidence about a candidate — their program, their council votes, their documented actions — and converts it into a comparable score a voter can read in 60 seconds on a phone.

#klimatapotrebuje's NRSR analysis proved the model works. Klima Kompas applies it to the local level with AI assistance and a clear MVP philosophy: get a working website with real scores in front of NGO stakeholders first, then automate incrementally.

Target Personas

Persona	Core Need
NGO Researcher	Tool that accepts manually entered evidence and produces a defensible score. Automation is enhancement, not prerequisite.
NGO Leadership	Working website to show stakeholders before committing to full data collection. The prototype is the pitch.
Slovak Voter	Badge answer in 60 seconds on mobile. No jargon. No login.
Media	Citable source with published methodology and a live URL.

2. SCORING MODEL — LOCKED

This scoring model is locked. It does not change between v5 and launch. Any revision requires a new methodology version, full re-scoring of all candidates, and public changelog on the methodology page.

2.1 Two-Pillar Structure

Pillar	Weight | Components
SLOVÁ — What they say	40% | Program analysis 15% + Questionnaire 15% + Social media 10% (Phase 2; = 0 in MVP)
SKUTKY — What they do	60% | Council votes 25% + Documented actions 35%

Formula: KLIMA_SCORE = (SLOVÁ × 0.40) + (SKUTKY × 0.60)

MVP SLOVÁ formula (social = null): SLOVÁ = (program × 0.50) + (questionnaire × 0.50)

Full SLOVÁ formula (Phase 2): SLOVÁ = (program × 0.375) + (questionnaire × 0.375) + (social × 0.25)

2.2 Badge Thresholds

Badge	Threshold | Label
🟢 Green	≥ 80% — Klimatický líder

🟡 Yellow	≥ 55% — Čiastočne angažovaný
🟠 Orange	≥ 30% — Minimálna angažovanosť
🔴 Red	< 30% — Klíma absentuje
⬜ Grey	Insufficient data — Nedostatok dát

Grey assignment rules: any required pillar is null AND not recoverable, OR overall confidence < 0.40, OR council vote evidence = 0 items AND documented actions evidence < 2 items (GREY_NEW_CANDIDATE). One grey badge visually. Sub-type shown as text in detail view only.

2.3 Normalisation — Absolute Caps (LOCKED)

All components use absolute min/max normalisation against predefined caps. Scores are stable — adding or removing a candidate never changes another candidate's score. Any candidate who reaches the cap earns 100%. Scores are floored at 0% and capped at 100% after normalisation. DB stores full float precision. Methodology page shows one decimal. Public display shows whole numbers.

Component	Min cap
Programs	-35.30
Questionnaire	0
Documented actions	-10
Council votes	-(n×2)

Normalisation formula (same for all):

normalised = min(100, max(0, (raw − min_cap) / (max_cap − min_cap) × 100))

Council votes example (8 climate votes in term, max_possible = 16):

- Kandidát A, raw +10 → min(100, max(0, (10+16)/32 × 100)) = 81%
- Kandidát B, raw -2 → min(100, max(0, (-2+16)/32 × 100)) = 44%

Program example:

- PS 2023, raw +11.01 → (11.01+35.30)/52.61 × 100 = 88%
- Smer 2023, raw -6.36 → (-6.36+35.30)/52.61 × 100 = 55%
- Venstre 2007, raw +17.31 → 100%
- Venstre 2015, raw -35.30 → 0%

2.4 SLOVÁ Component 1 — Program Analysis (Carter et al.)

Method: two independent raters classify every sentence in the candidate's election program as pro-climate, anti-climate, or neutral. A third person reconciles disagreements. Final score = pro_sentence_percentage − anti_sentence_percentage. This is the Carter et al. (2023, Party Politics) method applied verbatim.

AI adaptation: Claude API performs the initial classification pass with structured output. Two NGO reviewers verify a sample (minimum 20% of sentences per candidate). If agreement rate falls below 80%, full manual review is triggered.

Specificity weight (local adaptation): within pro-climate sentences, commitments are sub-scored 0 (vague aspiration), 1 (specific without timeframe), or 2 (specific with timeframe or measurable target). This sub-score adjusts the pro_sentence_percentage upward for high-specificity programs — acknowledging that a program with 3 specific commitments is more meaningful than one with 10 vague ones.

Local relevance filter: sentences referencing named local infrastructure, specific roads, rivers, neighbourhoods, or facilities score full weight. National-level responses to local questions score 0.5 weight.

> Normalisation caps (Carter et al. global range): Min cap: −35.30 (Venstre, Denmark 2015 — most anti-climate program in the dataset) Max cap: +17.31 (Venstre, Denmark 2007 — most pro-climate program in the dataset) Formula: min(100, max(0, (raw + 35.30) / 52.61 × 100)) 100% = matching the best climate program ever evaluated by this methodology across all countries. Any 2026 Slovak candidate scoring above +17.31 raw is capped at 100%.

2.5 SLOVÁ Component 2 — Questionnaire (NRSR +1/−1 Rubric)

8 questions targeting local competencies (transport, energy, green space, air quality, waste, water, climate strategy, EU funding). 6 open-ended, 2 binary (Áno/Nie/Neviem + optional justification).

Scoring: +1 per specific measure (clearly defined solution + purpose OR timeframe). −1 per anti-climate commitment. Local relevance filter applied as in program analysis.

If candidate does not respond: questionnaire_score = 0. This is a valid data point, not null. The candidate had the opportunity and did not engage. Does not alone trigger Grey badge.

> Normalisation caps (NRSR 2023 data): Min cap: 0 (non-respondent or no specific measures) Max cap: 54 points (Progresívne Slovensko, NRSR 2023 — highest score ever achieved) Formula: min(100, max(0, raw / 54 × 100)) 100% = matching the best questionnaire response ever submitted to #klimatapotrebuje. Any 2026 komunálne candidate scoring above 54 raw points is capped at 100%.

2.6 SLOVÁ Component 3 — Social Media (Phase 2)

3-axis classification per public post: environmental relevance (0–2), local specificity (0–2), concreteness (0–2). Post score = average of 3 axes. Aggregate with recency weight (posts in final 30 days before election count 0.5× to reduce last-minute greenwashing).

> Normalisation caps: TBD before Phase 2 implementation based on pilot data. A calibration run on existing candidate social media will establish realistic min/max anchors using the same absolute cap approach as all other components.

In MVP: social_score = null → SLOVÁ computed on program + questionnaire only (50/50 split). Every MVP scorecard displays: "Hodnotenie sociálnych sietí bude doplnené v ďalšej fáze."

2.7 SKUTKY Component 1 — Council Votes (25% of total)

Source: zastupiteľstvo and VUC plenary minutes, 2022–2026 term only. Verified via session date + agenda item number + URL.

Vote	Raw points
FOR a pro-climate resolution	+2
AGAINST a pro-climate resolution	-2
FOR an anti-climate measure	-2
AGAINST an anti-climate measure	+2
Abstained or absent	0 — neutral

> Normalisation caps (jurisdictional — varies per candidate): Min cap: $-(n \times 2)$ where n = number of climate-relevant votes available in this candidate's council term Max cap: $+(n \times 2)$ (same n) Formula: $\min(100, \max(0, (raw + n\times2) / (n\times4) \times 100))$ 100% = voted correctly on every single climate-relevant vote in the council term. 0% = voted incorrectly on every single climate-relevant vote. 50% = perfectly neutral (equal for and against, or entirely absent). The value of n is documented per jurisdiction on the methodology page.

New candidates: no council history → votes component = null. Does not alone trigger Grey. Evidence Density indicator on public card shows data basis.

2.8 SKUTKY Component 2 — Documented Actions (35% of total)

Covers all contexts: in office, in prior public roles, in civic or professional life. The only filter is: documented (URL required) + directly attributed to candidate by name + measurable output or clear position taken.

Action	Points
Climate strategy or SEAP adopted under candidate's leadership	+5
Joined Climate Alliance or Covenant of Mayors	+4
Environmental project initiated, documented budget ≥ €10,000	+3
Environmental policy submitted or co-submitted	+2
Led documented environmental project (civic, NGO, professional)	+2
Published environmental research or policy contribution	+2
Specific public commitment with timeframe or measurable target	+1
Environmental award, grant, or formal recognition	+1
Approved polluting project without environmental mitigation	-2
Reversed or blocked approved environmental measure	-2

Documented anti-environmental professional history	-2

Normalisation caps (designed methodology caps): Min cap: −10 (maximum theoretical negative: e.g. 5 polluting projects approved + blocked measure) Max cap: +15 (achieving strategy +5, covenant +4, project +3, policy +2, commitment +1) Formula: min(100, max(0, (raw + 10) / 25 × 100)) 100% = hitting the methodology's maximum designed score — requires documented major strategic leadership plus several concrete initiatives. 0% of the component = scoring −10 or below. 40% = score of 0 (no documented environmental actions, positive or negative).

2.9 Climate Relevance Framework

Every evidence item — from any source, entered by AI or human — must be tagged with a climate relevance tier before it contributes to scoring.

The tier is determined by HOW the candidate or council framed the action, not by the topic category of the action. The same vote on public transport can be Tier 1, Tier 2, or Tier 3 depending on whether the environmental connection is stated in the source or inferred by the reviewer. This prevents topic-based bias — public transport is not inherently "less direct" than EV chargers. What matters is whether the decision was made and framed as an environmental decision.

Tier	Definition | Classification Rule | Reviewer Action
Tier 1 — Explicit	The environmental connection is stated in the source document, vote proposal, or candidate's own statement. Keywords such as emisie, klíma, životné prostredie, or a measurable environmental target appear in the source itself. Examples: bus network expansion proposal that cites emission reduction targets; EV charger installation with documented CO2 savings; cycling infrastructure vote referencing air quality data; park budget framed as urban heat island mitigation. Rule: source document says it — reviewer cites the exact passage.
Tier 2 — Implicit	The environmental connection is real and demonstrable, but the source document does not frame it in environmental terms. The reviewer makes the connection and must document it. Examples: bus network expansion with no environmental framing in the proposal but with clear modal shift impact; cycling infrastructure voted on purely as transport policy; waste collection contract with no environmental language but clear material recovery impact. Rule: reviewer makes the link — must document why this action has a climate or environmental effect in citation_text.
Tier 3 — Excluded	No credible environmental connection — either not present or too indirect to document honestly. General economic, social, or infrastructure decisions with no environmental dimension stated or demonstrable. Examples: road widening without environmental framing; building permits without energy efficiency standards; school budgets without environmental programmes; general economic development votes. Rule: not scored. Stored in DB for transparency but excluded from all score computations. Requires secondary reviewer approval to reclassify to Tier 2.

Worked Example — Public Transport vs EV Chargers

This example illustrates why topic does not determine tier:

Scenario	Tier | Reason
Vote on regional bus network expansion. Proposal document cites: "Rozšírenie	Tier 1 | Environmental connection explicit in source — CO2 target stated

MHD zníži emisie CO2 o 12% v kraji."	
Vote on regional bus network expansion. Proposal framed as economic connectivity only. No environmental language.	Tier 2 | Reviewer documents: "Modal shift from private cars to public transport demonstrably reduces regional emissions"
Vote on installing 20 EV chargers at a shopping centre, framed as economic development.	Tier 2 | EV charging reduces tailpipe emissions; reviewer must document the link
Vote on EV chargers with documented emission reduction target in the proposal.	Tier 1 | Environmental connection explicit in source
Vote on road resurfacing. No environmental framing. No EV or cycling dimension.	Tier 3 | No credible environmental connection; excluded from scoring
Vote on road widening with dedicated cycling lane added and air quality report cited.	Tier 1 | Environmental connection explicit — air quality report cited in source

This framework is published verbatim on the public methodology page. It is the primary defence against "why is this vote considered climate-relevant?" challenges. Every Tier 1 item has a source citation. Every Tier 2 item has a reviewer-written justification visible to the public.

2.10 Candidate State Model

State	Meaning
REGISTERED	Candidate seeded into DB. No data collection yet.
DATA_COLLECTION	Evidence being gathered. Questionnaire dispatched.
ANALYZED	AI or manual scoring complete. Awaiting human review.
IN_REVIEW	Assigned to NGO reviewer.
NEEDS_REVISION	Reviewer flagged issues. Returned for re-analysis or additional evidence.
APPROVED	Score locked. Ready to publish.
PUBLISHED	Live on public site. Score versioned.

3. MVP DEFINITION

> MVP = minimum needed to put real scores on a real website for real voters. Nothing else. The website is the prototype. Show it to NGO stakeholders. Get organisational buy-in. Then build automation on top.

MVP Includes	Phase 2
Public website: badges, pillar bars, citations, methodology page	Candidate comparison
Scoring engine with locked formula	In-platform appeals system
Manual evidence entry (single item)	Evidence density indicator display

Review and approval dashboard	Audit history UI
Questionnaire submission acceptance (no management)	OG image generation
AI program + questionnaire analysis (CAP-02)	Bulk CSV/JSON import
Climate Relevance Framework enforced in entry form	Questionnaire reminders + state machine
	Social media agent (CAP-04)
	CAP-03 full automation

Appeals are handled manually outside the system in MVP. Questionnaire dispatch and reminders managed by NGO via email in MVP. OG images produced by NGO manually and posted on social media.

Capability Priority

Priority	Capabilities
P0 — Build first, in parallel	CAP-07 Scoring Engine + CAP-01 Public Website + CAP-06 Review Dashboard + CAP-08 Manual Evidence Entry
P1 — Build next	CAP-02 AI Program & Questionnaire Analysis + CAP-05-LITE Questionnaire Submission
P2 — MVP pilot only	CAP-03 Voting Record Parser (2–3 clean jurisdictions)
Phase 2	CAP-05-FULL Questionnaire Management + CAP-04 Social Media + comparison + appeals UI + evidence density + OG images + bulk import

4. DELIVERY TIMELINE

Phase	Scope | Target
Phase 0 — Prototype (now, ~1 week)	Fork repo. Supabase schema. 3 fake candidates with fake scores flowing end-to-end through formula to badge on website. Show to NGO stakeholders.
Phase 1 — MVP Build (~3–4 weeks)	All P0 and P1 capabilities operational. 16 real candidate shells live with Grey badges. Questionnaire links generated and sent manually.
Phase 2 — Data Collection (June–July, ~6 weeks)	Programs analyzed. SKUTKY manually entered for incumbents. Questionnaire responses ingested. CAP-03 piloted. First real scores appearing.
Phase 3 — Review Sprint (August, ~3 weeks)	All 16 candidates through review queue. Methodology page complete. All cards approved and published.
Phase 4 — Hard Launch (September)	PR, media, social. NGO posts OG images manually. Target: 50,000+ visitors before election day.
Phase 5 — Post-MVP (Nov 2026+)	Full automation, comparison, evidence density, appeals UI, city expansion.

Build timeline: 2–3 weeks coding + 4–6 weeks stabilisation (data reality, edge cases, reviewer UX iteration) = 6–9 weeks total to launch-ready.

5. LEGAL FRAMEWORK

- Zákon č. 181/2014 Z.z. o volebnej kampani: factual analysis permitted; endorsement prohibited.
- Every scorecard: "Toto hodnotenie nie je odporúčaním na hlasovanie."
- Every score point: cited with URL. No score without evidence.
- Climate Relevance Framework published — transparent classification rules.
- Questionnaire responses published verbatim, unedited.
- All candidates in each position evaluated — no selective coverage.

6. COST

Item	Cost
Claude Pro	$20/month
Lovable Pro (build phase, 2 months)	$50 total then downgrade
Antigravity	Free
Supabase	Free tier; upgrade at NGO discretion
Hosting	Lovable hosting (free) through development; production hosting TBD by NGO
Domain	~€15/year
Claude API (initial data run + monthly)	~$80 total over 6 months
Total infra (6 months)	~€300–380
NGO researcher time (operational)	~80–100 hours June–September 2026